

- [Home](#)
- [Archives](#)
- [Tags](#)
- [Categories](#)
- [About](#)

- [Home](#)
- [Archives](#)
- [Tags](#)
- [Categories](#)
- [About](#)

Passing file descriptors between processes via UNIX domain sockets

2022-01-04

[tutorials](#)

1209 words 6 min read

Table of Contents

- - [1. UNIX domain socket](#)
 - [2. sendmsg](#)
 - [3. recvmsg](#)

Linux provides a series of system calls that allow us to pass file descriptors between processes. Instead of simply passing the file descriptor, which is a 32-bit integer, we actually pass the file handle to the target process so that it can read and write to the file. Now suppose process B wants to send a file descriptor to process A. Let's see how to do that.

1. UNIX domain socket

To pass file descriptors, you need to establish inter-process communication. UNIX domain socket is a way to communicate between processes, it is similar to a normal socket, except that it does not use an ip address, but a socket file. We want to use it to send control messages and thus pass file descriptors. First we will have process A create a socket:

```
1 /*code of process A*/
2 int socket_fd = socket(AF_UNIX, SOCK_DGRAM, 0);
```

Note the arguments to the socket function. AF_UNIX specifies the protocol family as a UNIX domain socket. Here we use the datagram socket SOCK_DGRAM, which is similar to UDP and does not require a link to be established, but is sent directly through the address. Of course, you can also use SOCK_STREAM, which is similar to TCP. We won't list them here.

Next we bind the address:

```
1 /*code of process A*/
2 struct sockaddr_un un;
3 un.sun_family = AF_UNIX;
4 unlink("process_a");
5 strcpy(un.sun_path, "process_a");
6
7 if (bind(socket_fd, (struct sockaddr*)&un, sizeof(un)) < 0) {
8     printf("bind failed\n");
9     return 1;
10 }
```

Note that the address here is no longer an ip address, but a file path. Here we specify the address as process_a, and call bind to bind the address, which creates a socket file called process_a.

This way, other processes can send messages to this process through a special address like process_a. It is the same as sending UDP messages, except that struct sockaddr_un is used to define the address instead of struct sockaddr_in or struct sockaddr_in6. In addition, it is important to note that other processes can pass file descriptors to this process by sending control messages.

2. sendmsg

Linux provides a pair of system calls: sendmsg and recvmsg. Unlike our usual send and recv, they can be used to send or receive control messages in addition to regular data, which is the key to passing file descriptors; they can also be used to send or receive a discrete piece of data.

Let's look at the declarations of these two system calls

```
1 #include <sys/types.h>
2 #include <sys/socket.h>
3
4 ssize_t sendmsg(int sockfd, const struct msghdr *msg, int flags);
5 ssize_t recvmsg(int sockfd, struct msghdr *msg, int flags);
```

sendmsg and recvmsg use a structure struct msghdr to describe the data sent. The structure is defined as follows:

```
1 struct iovec {
2     void *iov_base;          /* Scatter/gather array items */
3     size_t iov_len;          /* Starting address */
4 }
```

```

3     size_t iov_len;          /* Number of bytes to transfer */
4 };
5
6 struct msghdr {
7     void      *msg_name;      /* optional address */
8     socklen_t  msg_namelen;    /* size of address */
9     struct iovec *msg_iov;     /* scatter/gather array */
10    size_t     msg_iovlen;     /* # elements in msg_iov */
11    void      *msg_control;     /* ancillary data, see below */
12    size_t     msg_controllen; /* ancillary data buffer len */
13    int        msg_flags;       /* flags on received message */
14 };

```

- `msg_name` Destination address. This is optional, if the protocol is connection-oriented, you don't need to specify an address; otherwise, you need to specify an address. This is similar to `send()` and `sendto()`.
- `msg_iov` the data to be sent. This is an array, the length of which is specified by `msg_iovlen`, and the elements of the array are a `struct iovec` structure that specifies the starting address (`iov_base`) and length (`iov_len`) of a contiguous piece of data. In other words, it can send more than one continuous piece of data; or, it can send a discontinuous piece of data.
- `msg_control` controls the message. This is what we are talking about today. We can't set it directly, we have to use a series of macros to set it.

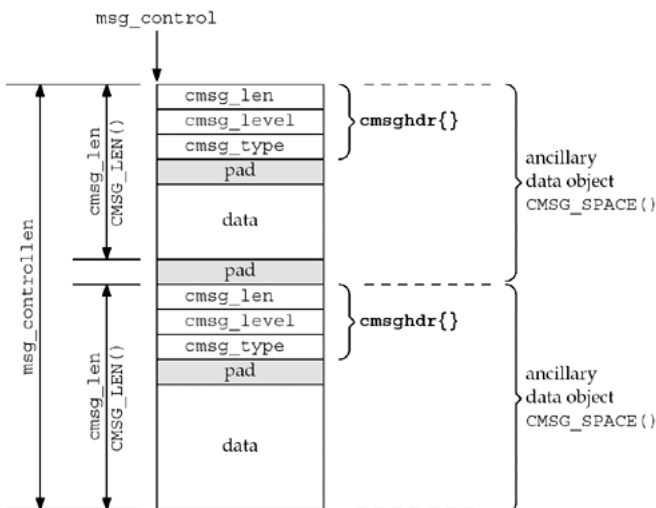
`msg_control` points to a sequence consisting of a `struct cmsghdr` structure and its additional data. The definition of `struct cmsghdr` is as follows:

```

1 struct cmsghdr {
2     socklen_t cmsg_len; /* data byte count, including header */
3     int       cmsg_level; /* originating protocol */
4     int       cmsg_type; /* protocol-specific type */
5 /* followed by
6    unsigned char cmsg_data[]; */
7 };

```

`struct cmsghdr` actually defines the header of the data, which should be followed by an unsigned `char` array that holds the actual data of the control information. This is often referred to as a variable-length structure. `msg_control` points to a sequence of such variable-length structures. The memory structure is shown in the figure below:



We need to use the following macros:

```

1 #include <sys/socket.h>
2
3 struct cmsghdr *CMSG_FIRSTHDR(struct msghdr *msgh);
4 size_t CMSG_SPACE(size_t length);
5 size_t CMSG_LEN(size_t length);
6 unsigned char *CMSG_DATA(struct cmsghdr *cmsg);

```

- `CMSG_FIRSTHDR()` returns the first element of the sequence pointed to by `msg_control`
- `CMSG_SPACE()` The length of the actual data of the incoming control message, returning the space required for the variable-length structure
- `CMSG_LEN()` The length of the actual data of the control message, and the length of the variable-length structure.
- `CMSG_DATA()` Returns the first address of the actual data that holds the control information.

Note the difference between `CMSG_SPACE()` and `CMSG_LEN()`: the former contains the length of the padding, which is the actual space occupied; the latter does not contain the length of the padding, and is used to assign the value to `cmsg_len`.

Next, let's have process B pass a file descriptor to process A. First set the address of process A, which is "process_a":

```

1 /*code of process B*/
2 struct sockaddr_un ad;
3 ad.sun_family = AF_UNIX;
4 strcpy(ad.sun_path, "process_a");

```

We only need to send control messages, not regular data, so we leave the regular data empty:

```

1 /*code of process B*/
2 struct iovec e = {NULL, 0};

```

Next, we allocate space for the control data, since we only pass one file descriptor, so the length is `sizeof(int)`:

```
1 /*code of process B*/
2 char cmsg[MSG_SPACE(sizeof(int))];
```

Then you can set the struct msghdr structure:

```
1 /*code of process B*/
2 struct msghdr m = {(void*)&ad, sizeof(ad), &e, 1, cmsg, sizeof(cmsg), 0};
```

Next we get struct cmsghdr and set it:

```
1 /*code of process B*/
2 struct cmsghdr *c = MSG_FIRSTHDR(&m);
3 c->cmsg_level = SOL_SOCKET;
4 c->cmsg_type = SCM_RIGHTS;
5 c->cmsg_len = MSG_LEN(sizeof(int));
6 *(int*)MSG_DATA(c) = cfd; // set file descriptor
```

Finally, just send it out:

```
1 /*code of process B*/
2 sendmsg(mfd, &m, 0)
```

3. recvmsg

Now we ask process A to receive the file descriptor passed to it. Here we call recvmsg .

```
1 /*code of process A*/
2 char buf[512];
3 struct iovec e = {buf, 512};
4 char cmsg[MSG_SPACE(sizeof(int))];
5 struct msghdr m = {NULL, 0, &e, 1, cmsg, sizeof(cmsg), 0};
6
7 int n = recvmsg(socket_fd, &m, 0);
8 printf("Receive: %d\n", n);
9
10 struct cmsghdr *c = MSG_FIRSTHDR(&m);
11 int cfd = *(int*)MSG_DATA(c); // receive file descriptor
```

Thus, process A receives the file descriptor from process B, and process B opens the file and can perform read and write operations on it.

[unix socket](#)

◀ [Enable Tab Search on your own blog](#) [Prev](#) [Promise and Deferred](#) [Next](#) ▶



